

1 Experiment Set Difference: Separable Convolutions (32 filters) v. Separable Convolutions (512 filters)

1.1 Average Treatment Effect (Per Experiment in Experiment Set)

| | |
|-----------------------------------|---|
| cpu_hours | -1.5477 +/- 0.0112163 (stderr) (p=7.03482e-13, Welch's t-test) |
| gpu_hours | -1.44208 +/- 0.00177115 (stderr) (p=1.6464e-16, Welch's t-test) |
| estimated_carbon_impact_kg | -0.130715 +/- 0.00874359 (stderr) (p=0.00174097, Welch's t-test) |
| total_power | -0.397482 +/- 0.0108522 (stderr) (p=0.000287525, Welch's t-test) |
| kw_hr_gpu | -0.222779 +/- 0.00051788 (stderr) (p=1.20092e-10, Welch's t-test) |
| kw_hr_cpu | -0.0287923 +/- 0.00722364 (stderr) (p=0.0236914, Welch's t-test) |
| exp_len_hours | -1.42663 +/- 0.00898631 (stderr) (p=9.55988e-10, Welch's t-test) |
| average_realtime_carbon_intensity | -21.8274 +/- 4.37705 (stderr) (p=0.000705002, Welch's t-test) |

1.2 Cumulative Difference (Per Experiment Set)

| | |
|-----------------------------------|----------------------|
| cpu_hours | -4.456080555555555 |
| gpu_hours | -4.225004845985649 |
| estimated_carbon_impact_kg | -0.38402716387715385 |
| total_power | -1.1664439532180533 |
| kw_hr_gpu | -0.6534909751171399 |
| kw_hr_cpu | -0.0847646914765648 |
| exp_len_hours | -4.111749448842473 |
| average_realtime_carbon_intensity | 559.1043773986441 |