

1 Experiment Set Difference: Conv (3-50 words) v. Transformer (3-50 words)

1.1 Average Treatment Effect (Per Experiment in Experiment Set)

| | |
|-----------------------------------|--|
| cpu_hours | -0.0108439 +/- 0.0102888 (stderr) (p=0.117613, Welch's t-test) |
| gpu_hours | -0.0429946 +/- 0.00379435 (stderr) (p=7.67862e-08, Welch's t-test) |
| estimated_carbon_impact_kg | -0.00127366 +/- 0.000613185 (stderr) (p=0.00851856, Welch's t-test) |
| total_power | -0.01346 +/- 0.00186468 (stderr) (p=1.51276e-06, Welch's t-test) |
| kw_hr_gpu | -0.00550392 +/- 0.000764456 (stderr) (p=9.08889e-06, Welch's t-test) |
| kw_hr_cpu | -0.00301504 +/- 0.000507361 (stderr) (p=5.75408e-05, Welch's t-test) |
| exp_len_hours | -0.0240631 +/- 0.0102563 (stderr) (p=0.00713085, Welch's t-test) |
| average_realtime_carbon_intensity | 62.6257 +/- 7.60088 (stderr) (p=0.00012757, Welch's t-test) |

1.2 Cumulative Difference (Per Experiment Set)

| | |
|-----------------------------------|-----------------------|
| cpu_hours | -0.0542194444444446 |
| gpu_hours | -0.2149732193209728 |
| estimated_carbon_impact_kg | -0.006368281068206236 |
| total_power | -0.06729978114492258 |
| kw_hr_gpu | -0.027519597314933267 |
| kw_hr_cpu | -0.015075200878055708 |
| exp_len_hours | -0.12031544718477472 |
| average_realtime_carbon_intensity | 313.128389651348 |